

Wake

a world model for organizational action

wake-web-zeta.vercel.app

WHO'S PRESENTING

Jack Gardner

Founder @ stealth startup · prev. founding tech lead @ 11x · ex-Amazon · ex-Williams F1

**You commit to a decision. The
world reacts.**

What if you could see how, before you commit?

THE QUESTION NO GOOD TOOL ANSWERS

THE PRODUCT LEAD

"If we ship this, will anyone care, or will it quietly enrage the people who love us?"

THE CEO

"If we announce this acquisition, what does our world look like in three months?"

THE ENGINEER

"If I take this idea to my manager, does it survive the room?"

Every one of these is answered by gut, by analogy, by a consultant substituting their gut for yours. None of it scales. None of it branches. None of it lets you watch the world you're about to act on.

THE ACTION

Notion is acquired by Microsoft.

What happens next?

Watch →

Two hundred minds, reasoning live.

A/B test the action before you take it.

WHY NO ONE BUILT THIS UNTIL NOW

Three things required, simultaneously.

01 · INTELLIGENCE PER NODE

Frontier-class reasoning

Each entity's behavior has to be worth reading and believing if you knew the person.

02 · OUTPUT SPEED

4× faster, ticks forward on stage

A 200-node cascade resolves in seconds. You cannot demo a simulation that runs overnight.

03 · COST AT SCALE

\$3.19 for 32 runs × 200 nodes

The Monte Carlo only works if the branches cost a few dollars, not thousands.

Frontier models had intelligence. Small models had speed. Neither had all three. **Gemini 3.5 Flash** is the first that does. *The cost basis is the product.*

THE GRAPH

The Notion world

208
NODES

346
EDGES

5
SEED ACTIONS

INTERNAL

47

C-suite · key managers ·
named engineers ·
employee archetypes by
team

EXTERNAL

38

Journalists · influencers ·
competitor leaders · VCs
who comment publicly

COMMUNITY

117

Productivity Twitter ·
skeptical HN · enterprise
IT buyer · switching
cohorts

PLATFORMS

6

Twitter · LinkedIn ·
Hacker News · Company
blog · Slack · Discord

THE SIMULATION KERNEL

Each node reasons, every tick.

```
// one Gemini 3.5 Flash call per node per tick

(node_state, inbox, world_clock)
→ (state_delta, outgoing_events, rationale)
```

INPUT

~600–1,000 tokens

Dossier cached as system prompt. Per-tick input is current state + inbox + world clock.

OUTPUT

~300–500 tokens

Structured JSON: state delta, typed outgoing events, a one-sentence first-person rationale.

COST

~\$0.005 / tick / node

Most input is cached. Lighter edges skip the LLM and use deterministic rules.

Wake.

sub-agents → every node, every edge

multi-step workflows → the cascade

long-horizon tasks at scale → the Monte Carlo

A world model for organizational action.

Wake.

What happens if Stripe acquires Plaid?

Build world →

✓ Researching the cast · Google Search

✓ Generating dossiers · Gemini 3.5 Flash

✓ Assembling world.json · runnable ✓

✓ Sizing to budget · 21 entities · ~\$0.016

✓ Writing edges & channels · 37 edges

✓ The cast · 21 entities

Same engine, any world, one sentence.

wake-web-zeta.vercel.app/genesis

HONEST FRAMING

Watch your decision before you make it.

Wake does not predict the future. It makes the consequence space of an action legible: the trajectories, the distribution, the pivotal variable.

World models are how embodied AI (robotics, self-driving) learns to reason about the physical world. Wake is that textual world model for **decisions**: one way to see any choice, any organization, and the impact it has on the world.

wake-web-zeta.vercel.app

Wake